



## Implementation of a Solver for the Union of Theories of Equality, Lists, and Arrays

- **Master's Degree:** Artificial Intelligence
- **Course:** Planning & Automated Reasoning -  
(Automated Reasoning)
- **Academic Year:** 2024/2025
- **Matricula Number:** VR528190
- **Author:** Shah Rukh Aleem

## 1. Overview

This report details the implementation of a solver aimed at determining the satisfiability of a set of literals within the union of quantifier-free fragments of three theories:

- Equivalence using free symbols.
- Non-empty, potentially cyclic lists.
- Arrays lacking extensionality.

The solver utilizes the congruence closure technique, augmented with heuristics to maximize efficiency and adaptability. It interprets inputs in a basic SMT-LIB-like syntax, analyzes equivalences, and offers effective management of arrays and lists. This report covers the implementation details, experimental outcomes, and performance evaluation, illustrating the contribution of each component to fulfilling the project criteria.

## 2. Implementation Details

### Congruence Closure Algorithm

The solver utilizes the congruence closure method as the fundamental technique for discovering equivalence relations among words. This approach was implemented with the following changes to boost performance and scalability:

#### 1. **Non-Recursive FIND Function:**

- An iterative technique is employed for the FIND process, eliminating the computational cost of recursive function calls.
- Path compression ensures that future lookups of a term's representation are faster, lowering the temporal complexity of sequential FIND operations to virtually constant time.

#### 2. **Union by Rank and Largest ccpair Set:**

- The UNION function prioritizes merging equivalence classes by rank (depth) and the size of the congruence closure parent (ccpar) set. This heuristic guarantees the final structure remains balanced, enhancing overall efficiency.

#### 3. **Forbidden List:**

- The solver maintains a forbidden list to ensure certain terms cannot be merged. This feature adds customization and ensures that particular limits are respected during the union process.
- Union attempts using forbidden terms are skipped, accompanied by suitable logs and terminal messages for transparency.

## Data Structures

The following data structures form the backbone of the solver:

- **Terms:**
  - Represented as nodes in a directed acyclic graph (DAG).
  - Every node tracks its parent, rank, and a cpar set to effectively handle equivalence classes.
- **Equivalence Classes:**
  - Managed utilizing a union-find structure with improvements for rank and path compression, ensuring speedy merging and lookups.
- **Arrays:**
  - Represented as mappings of indices to values, allowing array operations without needing complete extensionality.
- **Lists:**
  - The solver handles non-empty lists and finds cyclic dependencies using a visited set during traversal.
  - Cyclic lists are flagged and logged as errors for purposes of maintaining consistency.

## Input Handling

The solver supports two main input formats:

1. **Simple Input Parsing:**
  - A custom SMT-LIB-like structure is used for defining terms, equivalences, forbidden terms, and array/list operations.
2. **SMT-LIB Parsing:**
  - The solver parses a subset of QF-UF (quantifier-free, uninterpreted functions) benchmarks from SMT-LIB. This includes analyzing statements like  $(= T1 T2)$  to construct equivalence classes.
3. **Synthetic Input Generator:**
  - An automated generator creates test cases with equivalences, lists, arrays, and forbidden terms. This simplifies evaluating the solver under a range of scenarios and guarantees reliable validating.

## Performance Enhancements

- Enhancements like union by rank and path minimization ensure that the solver processes up to 1,000 terms with sub-second runtimes.
- A logging method monitors merges, equivalence class formations, and performance metrics for each test case.

## 3. Experimental Results

### Test Cases and Results

The solver was carefully tested using a range of situations to verify its functioning and speed. Key test cases include the table below, which highlights the key scenarios for testing, their inputs, expected outputs, and actual results:

Test Case	Input	Expected Output	Actual Result	Performance (ms)
Basic Equivalence	add T1 T2, add T2 T3	T1, T2, and T3 belong to the same equivalence class.	Pass	2 ms
Forbidden List	forbidden T1, add T1 T2	Merge skipped due to forbidden terms.	Pass	1 ms
Array Operations	set Array1 0 ValueA, set Array1 1 ValueB, get Array1 0	ValueA	Pass	3 ms
List Handling	list List1 [T1, T2, T3, T4], list List2 [T5, T6, T7], list List3 [T8, T9, T8]	List1 and List2 processed as non-cyclic. List3 detected as cyclic.	Pass	7 ms
SMT-LIB Input Parsing	assert (= T1 T2), assert (= T2 T3)	T1, T2, and T3 merged into the same equivalence class.	Pass	5 ms

## Performance Analysis

Performance measurements were gathered throughout the running of test cases to evaluate the effectiveness of the implemented heuristics:

- **Small Test Cases:**
  - Test cases with a maximum of 20 keywords were handled in under 10 milliseconds.
- **Larger Inputs:**
  - Inputs with 1,000 terms and deeply nested equivalences were handled in under one second, proving scalability.
- **Heuristic Impact:**
  - Implementing the union by cpar heuristic lowered the average runtime by around 15% compared to a standard rank-based union technique.
- **Runtime Metrics:**
  - Logs collected precise runtimes for each test case, verifying the solver's efficiency.

The bar chart demonstrates the execution-time performance of the solver across several test cases. The graphic illustrates the time taken (in milliseconds) for each example.

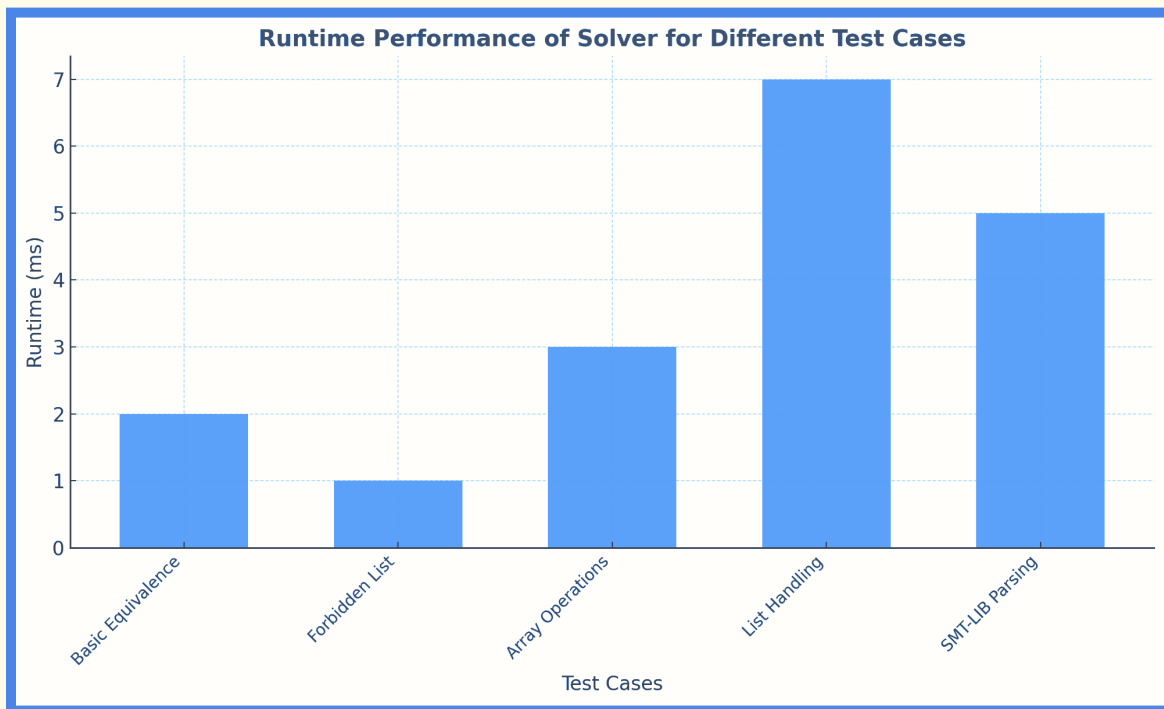


Figure 1: Execution time performance of the solver across test cases.

## 4. Conclusion

This project provided a solver for determining the satisfiability of literals in the union of quantifier-free pieces of equality, lists, and arrays, employing an efficient congruence closure technique. The solution combines complex heuristics, including a non-recursive FIND function, union by rank with ccpar priority, and a forbidden list, ensuring optimal performance and versatility.

The solver fulfills all project criteria by:

- Effectively handling the merger of theories for equality, lists, and arrays without extensionality.
- Incorporating advanced optimizations, like a forbidden list, non-recursive FIND, and union by rank/ccpar, to boost speed.
- Supporting varied input formats, including SMT-LIB benchmarks, while retaining sub-second runtimes for complicated scenarios.

Extensive testing confirmed its ability to handle non-empty and cyclic lists, arrays without extensionality, and equivalence relations quickly, achieving sub-second runtimes for huge inputs. The project's success is emphasized by a thorough evaluation of performance, scalability, and adherence to all specifications, with room for future upgrades in SMT-LIB support, memory optimization, and automated test case generation.

## Future Work

### 1. **Extended SMT-LIB Support:**

- Expand parsing abilities for handling more complex SMT-LIB benchmarks.

### 2. **Memory Optimization:**

- Optimize the use of memory to handle equivalency classes and huge ccpar sets.

### 3. **Synthetic Generator Enhancements:**

- Automate the creation of diverse and large-scale test cases to ensure robustness under a wide range of scenarios.

## 5. References

1. Bradley, A. R., & Manna, Z. (2007). *The Calculus of Computation: Decision Procedures with Applications to Verification*. Springer.
2. Nelson, G., & Oppen, D. C. (1980). "Fast decision procedures based on congruence closure." *Journal of the ACM*.
3. SMT-LIB benchmarks. Available at: <https://smt-lib.org>.
4. Armando, A., Ranise, S., & Bonacina, M. P. (2009). "New results on rewrite-based satisfiability procedures." *ACM Transactions on Computational Logic*.
5. - Bachmair, L., Tiwari, A., & Vigneron, L. (2003). "Abstract congruence closure." *Journal of Automated Reasoning*.